

Naga Renuka Kandi

Email: nagarenuka.k@gmail.com Mobile: +1 774-636-8864

Software Engineer | AI Evaluation | Python | Azure | AWS | REST APIs | LLM Systems

PROFESSIONAL SUMMARY

Software Engineer and AI Evaluation Analyst with 3+ years of experience building production backend systems, REST APIs, and cloud-native services on Microsoft Azure and AWS. Currently evaluating AI and LLM systems at LinkedIn, assessing response correctness, reasoning quality, and alignment with user intent across prompts and scenarios. Proven track record of shipping reliable, scalable software end to end and designing structured evaluation and automation pipelines for AI systems.

TECHNICAL SKILLS

AI & LLM Systems: AI/LLM Evaluation, Rubric-Based Assessment, Hallucination Analysis, Response Quality Analysis, REST APIs, Microservices, System Design, Backend Development, API Integration, Automation, Code Review, Technical Documentation

Languages: Python, C#, TypeScript, JavaScript, SQL, KQL, Bash

Backend & APIs: .NET, REST APIs, FastAPI, Flask, Microservices, Distributed Systems, Asynchronous Processing

Cloud & DevOps: Microsoft Azure (OneLake, Fabric, Monitor), AWS (EC2, S3, Lambda, SageMaker), Docker, Kubernetes, GitHub Actions, Jenkins, CI/CD

Databases & Data: SQL Server, PostgreSQL, Snowflake, DynamoDB, Redis, Query Optimization, ETL Pipelines, Data Modeling

Testing & Observability: Automated Testing (Cypress, pytest), KQL Telemetry, Azure Monitor, Logging, Tracing, Incident Response

PROFESSIONAL EXPERIENCE

LinkedIn

Jun 2026 - Present

AI Evaluation Analyst

- Evaluate AI and LLM generated responses for correctness, relevance, instruction following, reasoning quality, and overall response quality.
- Review model outputs to identify failure patterns and edge cases, documenting clear evaluation results and actionable feedback.
- Apply structured evaluation criteria and rubrics to assess model behavior consistently across different tasks and prompts.
- Analyze hallucinations, response relevance, personalization quality, reasoning quality, and whether responses correctly satisfy user intent.
- Perform structured testing of AI systems across different prompts and scenarios to identify quality issues and improvement opportunities.

Microsoft

Dec 2025 - Jun 2026

Software Engineer - Platform & AI Systems

- Built automated backend workflows on Azure OneLake, designing multi-step data pipelines that integrated structured telemetry queries to accelerate incident root cause analysis and reduce manual investigation time.
- Developed internal operational tooling by mapping incident scenarios to structured remediation content, validating system outputs against live platform behavior, and measuring accuracy to ensure operational reliability.

- Designed and iterated automation frameworks for incident resolution, structuring system inputs to improve triage consistency and surface actionable diagnostic insights from distributed Azure telemetry.
- Debugged distributed Azure microservices end-to-end using KQL queries in Azure Monitor, tracing cross-service API failures and surfacing degradation patterns through reliability dashboards.
- Developed .NET backend APIs and automated CI/CD pipelines using GitHub Actions, building reliable service layers and deployment workflows supporting Azure OneLake data services.

Handshake AI

May 2025 – Dec 2025

Software Engineer – AI Evaluation & Automation

- Designed and executed structured evaluation pipelines across 100+ production codebases, assessing software outputs against defined quality criteria and producing systematic assessments with written justifications to inform improvement decisions.
- Assessed code outputs across 50+ production-scale open-source repositories, validating correctness, completeness, and edge case coverage against defined golden solutions.
- Engineered input templates aligned with real GitHub PR objectives, iterating on structure to improve output quality across Python, JavaScript, TypeScript, and Java codebases.
- Authored realistic technical evaluation tasks spanning database management, data science, and system architecture domains, producing clear documentation with diagrams and code captures.
- Developed automated test suites including fail-to-pass and regression tests in Docker-based validation environments, enabling reproducible measurement of code quality and accuracy across heterogeneous codebases at scale.
- Performed structured output analysis, documenting why code submissions succeeded or failed against engineering benchmarks and providing actionable feedback to improve evaluation quality and coverage.

JPMorgan Chase & Co.

Feb 2022 – Aug 2023

Software Engineer I

- Built Python and SQL data pipelines powering React compliance dashboards for 200+ internal stakeholders, optimizing query performance by 30% through indexing strategies and data access layer redesign.
- Designed C# backend REST services with ASP.NET, integrating optimized SQL Server data layers and reducing page load times by 35% across finance operations dashboards.
- Scaled Cypress automated test coverage from 45% to 80%, catching 30+ pre-release defects and reducing post-deployment production incidents by 25% across distributed financial services.
- Automated Jenkins CI/CD pipelines with integrated testing gates and environment promotion, reducing release cycle time by 20% across QA and production environments.
- Supported AWS-hosted high-availability financial services (EC2, S3), triaging API failures, validating hotfixes, and maintaining uptime SLAs across enterprise deployments.

PROJECTS

promptLY – Privacy-First AI Writing Assistant

Published on Chrome Web Store | Chrome AI Challenge 2025

Built and published a Chrome extension using Gemini Nano for fully on-device LLM-powered text transformation with zero cloud dependency. Engineered prompt flows, preset transformations, side panel UI, context menu integration, and local storage architecture. Demonstrates end-to-end AI product engineering from architecture through public deployment.

Devil's Advocate Agent – Multi-Agent Reasoning System

HuggingFace Spaces | Groq | Llama 3.3 70B

Built a multi-agent debate system with role-based decomposition: Researcher retrieves counter-evidence, Devil's Advocate challenges the input stance, Defense reinforces it, and Judge evaluates both sides. Implemented iterative reasoning loops using Llama 3.3 70B via Groq API, deployed on HuggingFace Spaces with Gradio. Demonstrates agentic orchestration, prompt engineering, and structured multi-step LLM reasoning.

CERTIFICATIONS

- Agentic AI Certification - DeepLearning.AI (Andrew Ng) - 2026
- AWS Certified Solutions Architect - Associate
- AWS Certified Cloud Practitioner

EDUCATION

Master of Science - Analytics (Machine Learning Concentration) 2023 - 2025

Northeastern University, Boston, MA | GPA: 3.91/4.0

Bachelor of Engineering - Electrical and Electronics Engineering 2018 - 2022

VNR VJIT, Hyderabad, India